

# Fall 2023: CS5720

## Neural Networks & Deep Learning - ICP-7

**Name:** Manisha Lakkarsu

**Student Id:** 700746573

**Git link:**

**Video link:** [https://github.com/Mani543/Manisha\\_NNDL\\_ICP7.git](https://github.com/Mani543/Manisha_NNDL_ICP7.git)

[https://drive.google.com/file/d/12xBVJj3eHGERIQM\\_jAQhBjFEVfcIVUxL/view?usp=sharing](https://drive.google.com/file/d/12xBVJj3eHGERIQM_jAQhBjFEVfcIVUxL/view?usp=sharing)

**In class programming:**

1. Follow the instruction below and then report how the performance changed. (apply all at once)
  - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
  - Dropout layer at 20%.
  - Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
  - Max Pool layer with size 2×2.
  - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
  - Dropout layer at 20%.
  - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
  - Max Pool layer with size 2×2.
  - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
  - Dropout layer at 20%.
  - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
  - Max Pool layer with size 2×2.
  - Flatten layer.
  - Dropout layer at 20%.
  - Fully connected layer with 1024 units and a rectifier activation function.
  - Dropout layer at 20%.
  - Fully connected layer with 512 units and a rectifier activation function.
  - Dropout layer at 20%.
  - Fully connected output layer with 10 units and a Softmax activation function.

```

In [12]: import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.constraints import MaxNorm
from tensorflow.keras.utils import to_categorical
from keras.optimizers import SGD

# Fix random seed for reproducibility
np.random.seed(7)

# Load data
(a_train, b_train), (a_test, b_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
a_train = a_train.astype('float32') / 255.0
a_test = a_test.astype('float32') / 255.0

# One hot encode outputs
b_train = to_categorical(b_train)
b_test = to_categorical(b_test)
num_classes = b_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu'))
model.add(Dropout(0.2))

```

```

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 5
lRate = 0.01
decay_rate = lRate / epochs
sgd = SGD(learning_rate=lRate, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

```

```

# Fit the model
history = model.fit(a_train, b_train, validation_data=(a_test, b_test), epochs=epochs, batch_size=32)

# Evaluate the model
scores = model.evaluate(a_test, b_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

```

## Output:

Model: "sequential_3"		
Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 32, 32, 32)	896
dropout_18 (Dropout)	(None, 32, 32, 32)	0
conv2d_19 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_20 (Conv2D)	(None, 16, 16, 64)	18496
dropout_19 (Dropout)	(None, 16, 16, 64)	0
conv2d_21 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_22 (Conv2D)	(None, 8, 8, 128)	73856
dropout_20 (Dropout)	(None, 8, 8, 128)	0
conv2d_23 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_11 (MaxPooling2D)	(None, 4, 4, 128)	0
max_pooling2d_11 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_3 (Flatten)	(None, 2048)	0
dropout_21 (Dropout)	(None, 2048)	0
dense_9 (Dense)	(None, 1024)	2098176
dropout_22 (Dropout)	(None, 1024)	0
dense_10 (Dense)	(None, 512)	524800
dropout_23 (Dropout)	(None, 512)	0
dense_11 (Dense)	(None, 10)	5130
Total params: 2915114 (11.12 MB)		
Trainable params: 2915114 (11.12 MB)		
Non-trainable params: 0 (0.00 Byte)		
None		
Epoch 1/5		
1563/1563 [=====] - 216s 137ms/step - loss: 1.8654 - accuracy: 0.3102 - val_loss: 1.5189 - val_accuracy: 0.4385		
Epoch 2/5		
1563/1563 [=====] - 221s 141ms/step - loss: 1.4289 - accuracy: 0.4789 - val_loss: 1.2839 - val_accuracy: 0.5338		
Epoch 3/5		

```
Epoch 1/5
1563/1563 [=====] - 216s 137ms/step - loss: 1.8654 - accuracy: 0.3102 - val_loss: 1.5189 - val_accu
racy: 0.4385
Epoch 2/5
1563/1563 [=====] - 221s 141ms/step - loss: 1.4289 - accuracy: 0.4789 - val_loss: 1.2839 - val_accu
racy: 0.5338
Epoch 3/5
1563/1563 [=====] - 245s 157ms/step - loss: 1.2073 - accuracy: 0.5655 - val_loss: 1.0837 - val_accu
racy: 0.6123
Epoch 4/5
1563/1563 [=====] - 250s 160ms/step - loss: 1.0507 - accuracy: 0.6275 - val_loss: 1.0062 - val_accu
racy: 0.6406
Epoch 5/5
1563/1563 [=====] - 234s 150ms/step - loss: 0.9395 - accuracy: 0.6662 - val_loss: 0.9356 - val_accu
racy: 0.6768
Accuracy: 67.68%
```

- Did the performance change?
  - With the addition of more layers and feature maps, the model's performance is likely to improve, but the complexity and training time of the model will also rise. The new model architecture described in the instruction has more feature maps and new layers, which could increase the model's accuracy.

### In class programming:

2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly Image Classification on the handwritten digits data set (mnist)

```
In [13]: # Predict the first 4 images of the test data
         predictions = model.predict(a_test[:4])

         # Convert the predictions to class labels
         predicted_labels = np.argmax(predictions, axis=1)

         # Convert the actual labels to class labels
         actual_labels = np.argmax(b_test[:4], axis=1)

         # Print the predicted and actual labels for the first 4 images
         print("Predicted labels:", predicted_labels)
         print("Actual labels:", actual_labels)

1/1 [=====] - 0s 450ms/step
Predicted labels: [3 1 8 0]
Actual labels: [3 8 8 0]
```

## In class programming:

### 3. Visualize Loss and Accuracy using the history object.

```
In [15]: import matplotlib.pyplot as plt

# Plot the training and validation Loss
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()

# Plot the training and validation accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```

## Output:

